

NIKHIL GADE

Charlotte, NC | nikhilgade.me@gmail.com | [linkedin.com/in/nikhil--gade](https://www.linkedin.com/in/nikhil--gade) | github.com/3nikl

PROFESSIONAL EXPERIENCE

AI Engineer | Hanu Finance Corp, Charlotte, NC

Feb 2026 - Present

- Built and shipped multi-agent LLM systems on AWS for a multi-tenant AI contact center used by 800+ people daily, covering chat and an Amazon Connect voice agent with licensure-aware routing
- Dropped p95 latency from 4.2s to 1.1s via Redis semantic caching and retry/fallback paths for timeouts
- Designed LangGraph agent workflows with tool calling, conversation memory, and RAG over internal knowledge, gated by an LLM-as-judge eval suite that blocked regressions across 15+ prompt releases
- Reduced LLM spend 40% by routing simple requests to cheaper models, with guardrails on every action

Software Engineer Co-op | Hanu Finance Corp (Client: Wells Fargo), Charlotte, NC

Jan 2025 - Jul 2025

- Owned Java Spring Boot telemetry ingestion for 2.3M financial records daily with idempotent writes
- Supported an event-driven AWS Lambda/DynamoDB pipeline with REST APIs, cutting fraud detection 65%
- Contributed to an Azure OpenAI RAG chatbot over 450+ cards and merchant services documents for 2,500+ operations and care agents, tracking latency, token cost, and failures in Application Insights

Software Engineer | AI Variant, India

Apr 2021 - Aug 2023

- Improved reporting query performance 60% by normalizing schemas and adding covering indexes
- Rebuilt Python ETL pipelines on 8M records/run, cutting runtime from 4 hours to 34 minutes
- Built a private docs chatbot on local open-source LLMs and FAISS, so data never left the network

TECHNICAL SKILLS

Languages: Python, SQL, Java, JavaScript, TypeScript

Backend and Data: FastAPI, Spring Boot, Apache Kafka, Apache Flink, Airflow, PostgreSQL, MySQL, DynamoDB, Redis, Microservices, Distributed Systems

Reliability and Testing: Deterministic Simulation Testing, Fault Injection, Idempotency, Exactly-Once Semantics, pytest

Cloud and Tools: AWS, Azure, Docker, Kubernetes, CI/CD, MLflow, Grafana, Application Insights, Claude Code, Terraform, Git, Linux

AI and ML: LLMs (OpenAI, Anthropic, Azure OpenAI), LangChain, LangGraph, Multi-Agent Systems, RAG, Vector Search (Qdrant, Pinecone, FAISS), Prompt Engineering, LLM Evaluation, Guardrails, Semantic Caching, QLoRA, scikit-learn, XGBoost

Frontend: React, Next.js, Tailwind CSS

PROJECTS

PagedOut: Autonomous Incident Remediation and Simulation Testing | github.com/3nikl/pagedout

Independent Project

Python, LangGraph, Ollama (phi3), Apache Kafka, Apache Flink, Qdrant, RocksDB, FastAPI, Docker

- Engineered Groundhog, a deterministic simulator with a virtual clock and lossy network that ran 600K failure scenarios at 117,000x compression, catching 2 double-execution bugs a 200-run test suite missed
- Built a Kafka and Flink pipeline sustaining 32K events/sec at 68ms p99, and lifted runbook Recall@3 from 67.7% to 90.6% with hybrid dense and BM25 retrieval
- Designed a LangGraph agent graph that triages incidents, retrieves runbooks, and plans fixes, escalating cases under 0.5 confidence to a human and keeping action risk in a registry instead of the model

NephroSense AI: ML Risk Prediction Platform (Client: Kingdom Capital)

Graduate Capstone

Python, scikit-learn, XGBoost, SHAP, FastAPI, PostgreSQL, Docker, Azure

- Dropped p95 query latency from 820ms to 210ms by collapsing 14 N+1 lookups into 2 batched queries
- Led a 6-person team, working with client stakeholders and running weekly demos across 4 milestones
- Built an end-to-end ML pipeline on 3,200+ records (feature engineering, training, inference), benchmarking Logistic Regression, Random Forest, and XGBoost to 88% ROC-AUC and 94% recall with SHAP

EDUCATION

Master's in Informatics | Northeastern University, Boston, MA

Sep 2023 - Dec 2025

GPA: 3.9/4.0